

Sparse-Autoencoder Dissection of Emotion Features in Frozen Speech Encoders

Nima Kelidari

USC ICT, Intelligent Human Perception Lab
Los Angeles, CA, USA
kelidari@usc.edu

Chaitanya Parwatkar

University of Southern California
Los Angeles, CA, USA

Minoo Ahmadi

USC ICT, Intelligent Human Perception Lab
Los Angeles, CA, USA

Xiangxu Lin

University of Southern California
Los Angeles, CA, USA

Abstract

Linear probes show emotion is decodable in the middle-to-late layers of frozen speech transformers, but not which internal features carry it, nor whether the signal is acoustic or lexical. We train TopK Sparse Autoencoders (SAEs) on *per-frame* activations of six frozen large encoders and dissect their emotion features on the lexicon-controlled CREMA-D corpus (91 actors $\times$ 12 fixed sentences $\times$ 6 emotions). A confound-controlled selection—a feature’s firing must be more informative about emotion than about the spoken sentence or the speaker—isolates genuine emotion features from phonetic and speaker-driven ones. We find that (i) emotion *disentangles* into clean monosemantic features earlier (L2–L7) than where it is best linearly *decoded* (L10); (ii) all six encoders share an acoustic emotion-feature basis that generalizes across unseen speakers; and (iii) the recovered emotion signal is acoustic, not lexical. Decodability and disentanglement are thus distinct, and emotion in these encoders is carried by speaker-invariant acoustic features.

Keywords

speech emotion recognition, sparse autoencoders, mechanistic interpretability, probing, acoustic features

1 Introduction

Self-supervised speech encoders dominate speech emotion recognition, yet were trained on ASR-style objectives that encode linguistic structure. Probing reveals *where* emotion is readable, not *which* features encode it or whether the signal is acoustic or lexical. Sparse autoencoders (SAEs) decompose dense activations into sparse, near-monosemantic features [?], moving the analysis from decodability to mechanism. The fixed lexicon of CREMA-D [?] dissociates acoustic delivery from text by construction, so any recovered emotion signal there is necessarily acoustic.

2 Method

We re-extract *un-pooled* per-frame hidden states from six frozen large encoders (HuBERT [?], WavLM, wav2vec2, w2v-BERT, MERT, Whisper [?]); the project’s cached embeddings are mean-pooled, too few to train an over-complete SAE. CREMA-D layer 13 yields 940,273 frames. We train a TopK SAE ($d_{in}=1024$, $d_{sae}=8192$, $k=32$, AuxK dead-feature revival), reaching FVU = 0.070 (93% variance explained), with 4/8192 dead features.

Emotion features are scored by the mutual information (MI) between a feature’s firing and the emotion label, thresholded at 5σ over a label-permutation null. **Confound control is decisive:** a feature qualifies only if MI(emotion) exceeds both MI(sentence) and MI(actor), and it is monosemantic (≥ 0.5 activation mass on one emotion). Significance and monosemanticity alone flag 6407 significant / 394 monosemantic features, but $\sim 66\%$ carry more sentence (phonetic) or speaker information than emotion—the documented “SAE rediscovers phonemes, not affect” failure [?]. After control, 61 genuine emotion features remain (median monosemanticity 0.59), dominated by anger and sadness; fear and neutral drop to zero. Acoustic-vs-lexical labeling uses cross-validated ROC-AUC of a logistic classifier predicting firing from eGeMAPS [?] (acoustic) vs. sentence identity (lexical); R^2 is unusable as the pooled features are $\sim 96\%$ sparse.

3 Results

3.1 Depth: decode $\neq$ disentangle (Q i)

A 25-layer sweep separates two curves (Fig. 1, top-left). Linear probe accuracy peaks at L10 (0.739); the count of monosemantic emotion features peaks earlier at L7, and monosemanticity at L2. Deep layers (L17–24) retain decodability (~ 0.72) but lose disentanglement (24–47 features): emotion stays linearly readable while becoming distributed. *Emotion organizes into clean features before the depth at which it is best decoded.*

3.2 Generalization across paradigms and speakers (Q ii)

All six encoders carry emotion features at their best layer (HuBERT 61, wav2vec2 70, Whisper 79, WavLM 145, MERT 64, w2v-BERT 61), all anger/sadness-dominated. Matching features across encoders by per-utterance firing correlation, Whisper $\rightarrow$ SSL overlap (0.40) $\approx$ SSL $\rightarrow$ SSL overlap (0.39): on fixed-lexicon CREMA-D Whisper is *not* distinct, its known divergence concerning the lexical shortcut rather than acoustic emotion encoding. Under speaker-disjoint cross-validation over 91 actors, emotion accuracy is unchanged (random 0.551 vs. disjoint 0.551; leave-speaker-out drop ≈ 0) and each feature fires for $\sim 95\%$ of its emotion’s actors. Because selection excluded speaker-entangled features, the published MESD leave-one-speaker-out collapse (~ 0.30) is attributable to exactly those excluded features.




Figure 1: SAE dissection of frozen-encoder emotion features (HuBERT / CREMA-D). Top, left to right: depth decode-vs-disentangle (probe accuracy vs. # monosemantic emotion features); emotion features acoustic vs lexical (eGeMAPS vs. sentence firing-AUC); cross-encoder feature overlap. Bottom: emotion features per encoder; cross-speaker (leave-speaker-out) generalization; causal sufficiency of acoustic vs. lexical vs. random features.

Table 1: Key numbers (HuBERT / CREMA-D).

Quantity	Value
SAE FVU (variance explained)	0.070 (93%)
Dead features	4/8192
Emotion features (sig → mono → confound)	6407 → 394 → 61
Decode peak / disentangle peak layer	L10 / L7
Acoustic features (positive control)	60/61
Sufficiency acoustic / lexical / random	0.552/0.431/0.423
Cross-speaker leave-speaker-out drop	≈ 0

On lexicon-controlled data all encoders converge on a shared acoustic basis whose features survive speaker-disjoint evaluation; the Whisper-vs-SSL divergence is about exploiting a lexical shortcut when one exists [?], which CREMA-D cannot pose and which the deferred EMIS experiment targets. Utterance-pooled sparse features make necessity-ablation of small subsets near-null (the signal is distributed), so sufficiency is the informative test. The IEMOCAP real-transcript ablation, the SpeechCraft cross-lingual arm, and the EMIS audio-vs-text separation are deferred (data unavailable on our cluster) with runnable stubs.

3.3 Acoustic, not lexical (Q iii)

The positive control (fixed lexicon $\Rightarrow$ must be acoustic) labels 60/61 emotion features acoustic (eGeMAPS firing-AUC 0.83 vs. sentence AUC 0.66). Causally (Fig. 1, bottom-right), emotion decodes from the 61 acoustic features at 0.552, vs. 0.431 from lexical and 0.423 from random features (6-class; full reconstruction 0.73). The recovered emotion signal is acoustic.

4 Discussion and Limitations

Decodability and disentanglement are distinct axes: a probe reads emotion best where the representation is linearly separable, but the encoder packs emotion into clean monosemantic features earlier.